

COGS Negotiation Agent — How It Works (End to End)

A Databricks-native, agentic application that helps **Kroger category negotiators** prepare for and run supplier COGS (cost-of-goods-sold) negotiations across the beverage category (PepsiCo, Coca-Cola, Keurig Dr Pepper).

It is built entirely on real Databricks primitives — Unity Catalog, Metric Views, Genie, Vector Search, Mosaic AI Model Serving, Lakebase, AI/BI dashboards, and a Databricks App — with a **pluggable LLM layer** (Mosaic AI Gateway ↔ LiteLLM).

1. What it does

A negotiator can:

Capability	What it produces	Powered by
Negotiator (AI)	Ask anything in plain English; it routes to the right tool	Multi-Agent Supervisor
Supplier Scorecard	Live metrics table + conversational Q&A + an AI/BI dashboard	Genie + AI/BI
Negotiation Brief	Board-ready talking points grounded in the supplier's contract	Knowledge Assistant (RAG)
Fact-Pack Deck	A data-driven negotiation deck (KPIs, narrative, asks)	Deck builder (LLM)
Rehearsal Room	Role-play practice against the supplier's account manager	Vendor Rehearsal agent

2. Architecture (three layers)

```
flowchart TB
    subgraph L1["Layer 1 · Experience – Databricks App (React + FastAPI)"]
        UI["Negotiator chat · Deck Hub · Scorecard · Brief · Deck · Rehearsal"]
    end

    subgraph L2["Layer 2 · Runtime"]
        SUP["Multi-Agent Supervisor (LangChain router)"]
        LLM["Pluggable LLM factory get_llm()"]
        AGENTS["Sub-agents: Brief (RAG) · Deck · Rehearsal · Scorecard"]
        SERVED["Agent-as-model (Mosaic AI Model Serving)"]
    end

    subgraph L3["Layer 3 · Data & Governance – Unity Catalog"]
        DELTA["Delta tables + Metric View"]
        GENIE["Genie space (NL→SQL)"]
        VS["Vector Search index over contract MSAs"]
        VOL["Volume (contract docs)"]
        LB["Lakebase (Postgres OLTP)"]
        BI["AI/BI dashboard"]
    end

    PROV["Foundation Models: Claude Sonnet 4.5 (FMAPI) · or LiteLLM proxy"]

    UI --> SUP --> AGENTS --> LLM --> PROV
    AGENTS --> GENIE --> DELTA
    AGENTS --> VS --> VOL
    UI --> GENIE
    UI --> LB
    UI --> BI --> DELTA
    SERVED --> LLM
```

- **Layer 1 — Experience:** the Databricks App users interact with.
 - **Layer 2 — Runtime:** the agent brains — a supervisor that routes to sub-agents, all calling the pluggable LLM factory; also packaged as a served model.
 - **Layer 3 — Data & Governance:** everything in Unity Catalog — the single governance boundary.
-

3. Repository layout

```
cogs-negotiation-agent/
├─ app.py                # FastAPI entry – mounts API + serves the React SPA; opens
├─ app.yaml              # Databricks App config + all env vars (LLM_PROVIDER, GENIE)
├─ pyproject.toml / uv.lock # Python deps (app runtime)
├─ agent_model.py        # The agent packaged as an MLflow ChatAgent (for serving)
├─ deploy_agent.py       # Log → register (UC) → deploy the served model
├─ data_foundation.py    # Builds Delta tables + Metric View
├─ build_genie.py         # Authors the Genie space
├─ build_knowledge.py     # Builds contract Volume + chunks + Vector Search index
├─ build_dashboard.py    # Builds/updates + publishes the AI/BI dashboard
├─ eval_agent.py         # Quality Loop (offline): LLM-judge eval over the agent
├─ monitor_agent.py      # Quality Loop (online): scheduled judges on live traces
├─ server/
│   ├─ config.py         # Dual/triple-mode auth (local CLI / App SP / model-serving)
│   ├─ llm.py            # ★ Pluggable LLM factory – get_llm(), Mosaic ↔ LiteLLM
│   ├─ data.py           # Synthetic source-of-truth data (mirrored into UC)
│   ├─ agents.py         # Brief (RAG) / Deck / Rehearsal LangChain chains
│   ├─ supervisor.py     # Multi-Agent Supervisor (LLM router)
│   ├─ genie.py          # Genie Conversation API client
│   ├─ knowledge.py      # Vector Search retrieval (RAG)
│   ├─ state.py          # Lakebase (Postgres) persistence + in-memory fallback
│   └─ routes/           # FastAPI routers: catalog, agentic, genie_routes, supervi
└─ frontend/            # React + Vite + TS + Tailwind
    └─ src/pages/        # Hub, Negotiator, Scorecard, Brief, Deck, Rehearsal
```

4. Component-by-component

4.1 The pluggable LLM factory — `server/llm.py` ★

The heart of the design. Every chain/agent calls `get_llm()` and receives a LangChain `BaseChatModel`; **none of them know which provider is live**.

- `LLM_PROVIDER=mosaic` (default) → `ChatOpenAI` pointed at the **governed Databricks path**. On this Azure workspace the dedicated `ai-gateway.*` subdomain isn't available, so it falls back to the workspace `/serving-endpoints` OpenAI-compatible route (usage still records in `system.serving.endpoint_usage`; Gateway config attaches at the endpoint).
- `LLM_PROVIDER=litellm` → the same `ChatOpenAI` pointed at a **LiteLLM proxy** (`LITELLM_BASE_URL`). LiteLLM can route upstream *through* Mosaic to keep one governance boundary, or go direct.

- Both providers are OpenAI wire-compatible, so the switch is **one env var, zero code change**. A third `databricks` mode (native `ChatDatabricks`) exists behind a flag.

Switching is config-only — change `LLM_PROVIDER` in `app.yaml` (app) or the serving endpoint env (served model) and redeploy.

4.2 Auth — `server/config.py`

Detects the runtime and authenticates accordingly:

- **Local dev** → Databricks CLI profile (`DATABRICKS_PROFILE`).
- **Databricks App** → auto-injected service-principal credentials.
- **Model serving** → automatic-auth identity (scoped to declared resources).

4.3 Data foundation (Layer 3) — `data_foundation.py` → `kroger_demo.cogs`

- `supplier_scorecard` (Delta) — one row per supplier with spend, COGS/unit, landed-cost index, YoY changes, trade funds, fill rate, OTIF, rebate tier, contract expiry.
- `regional_performance` (Delta) — dollar sales, YoY, units, store counts by US Census division.
- `landed_cost_metrics` (**Metric View**) — the *certified* definitions of Total Spend, Avg COGS/Unit, Blended Landed Cost Index, Blended COGS Inflation, etc. Queried with `MEASURE(...)`. This is the single source of truth that Genie and the dashboard agree on.

The numbers originate in `server/data.py` so the app and UC stay in sync; in a real deployment these tables come from ingestion + a certified metric layer.

4.4 Genie space — `build_genie.py`

A Genie space (`id 01f1642219cf135cb84f7e0dbc8d6957`) over the metric view + tables, with instructions encoding the domain rules (use `MEASURE()`, landed-cost-index baseline = 100, positive YoY = cost inflation), example SQL, and benchmark questions. Turns plain-English questions into governed SQL.

4.5 Knowledge Assistant (RAG) — `build_knowledge.py` + `server/knowledge.py`

- Per-supplier **Master Supply Agreements** (6 clauses each: Pricing, COGS Adjustment, Rebates, Trade Funds, Service Levels, Term/Termination) written to a UC **Volume** `/Volumes/kroger_demo/cogs/contracts/`.
- Clause-level **chunks** in Delta `contract_chunks` (Change Data Feed on).

- A **Delta-Sync Vector Search index** `contract_chunks_index` with **managed embeddings** (`databricks-gte-large-en`) on the existing `kroger-recipe-search` endpoint.
- `knowledge.retrieve()` pulls the top-k clauses *filtered to the supplier*; the Brief agent injects them and **cites the section names**.

4.6 Sub-agents — `server/agents.py`

LangChain chains, all on `get_llm()` :

- `build_negotiation_brief` — RAG-grounded brief (retrieves contract clauses → prompt → brief that cites clauses + numbers).
- `build_fact_pack` — returns a structured JSON deck (title, hypothesis, KPIs, sections, asks) the React deck viewer renders.
- `rehearse_turn` — role-plays the supplier's key account manager, in character.

4.7 Multi-Agent Supervisor — `server/supervisor.py`

A pure-LangChain LLM **router** (no LangGraph). Given any message it classifies intent → `{scorecard, brief, deck, rehearse, chat}` and extracts the supplier, then dispatches:

- `scorecard / chat` → **Genie**
- `brief` → RAG brief · `deck` → deck builder · `rehearse` → rehearsal agent

Returns a unified envelope the UI renders by `route` (text / table+SQL / deck).

4.8 Agent-as-model — `agent_model.py` + `deploy_agent.py`

The same logic packaged as an **MLflow ChatAgent** , registered in Unity Catalog (`kroger_demo.models.cogs_negotiation_agent`) and deployed to a **Mosaic AI Model Serving** endpoint (`agents_kroger_demo-models-cogs_negotiation_agent`). This gives the agent its own queryable REST endpoint, inference tables, evaluation, and monitoring.

- Driven by `custom_inputs.task = brief | deck | rehearse | chat` .
- Declares the FM endpoint, embedding endpoint, Vector Search index, Genie space, and SQL warehouse as **serving resources** (automatic auth).
- **The LLM switch is preserved** here too — set `LLM_PROVIDER` as the endpoint env.
- **Design note:** the served model is **self-contained** — `brief` (RAG), `deck` , `rehearse` , and a grounded `chat` over bundled facts. Genie-backed NL→SQL routing lives in the **app**, because a served model runs Genie under an automatic-auth identity that can't execute the downstream SQL.

4.9 Lakebase persistence — `server/state.py`

Saved briefs/decks/rehearsals + the work queue persist to **Lakebase (Postgres OLTP)** instance `cogs-lakebase`, table `agent_artifacts(id, kind, supplier_key, payload jsonb, created_at)`. Tokens are minted fresh per connection via `w.database.generate_database_credential(...)` (no background refresh). If Lakebase isn't attached, it falls back to an in-memory store so the app still runs.

4.10 AI/BI dashboard — `build_dashboard.py`

A published **Lakeview dashboard** (`id 01f1643845191b049c9a3acf7531495b`) over `kroger_demo.cogs`: KPI counters, supplier comparison bars, regional bars, a detail table, and a **global-filters page** (Supplier / Rebate Tier / Region). It's **embedded** in the Scorecard page via iframe (workspace embedding policy set to `ALLOW_APPROVED_DOMAINS` for the app's domain) with an "open full dashboard" link.

4.11 Quality Loop — `eval_agent.py` + `monitor_agent.py`

The deployed agent is observed and judged, not just shipped:

- **Inference tables (automatic).** `agents.deploy()` enables request/response + trace logging to `kroger_demo.models.cogs_negotiation_agent_payload`. Every call to the served agent is captured.
- **Offline eval with LLM judges** (`eval_agent.py`). An eval dataset (brief / scorecard / chat / rehearse cases with `expected_facts`) is run through the served endpoint and scored by Mosaic AI / MLflow judges: **Correctness**, **RelevanceToQuery**, **Safety**, plus two custom **Guidelines** judges (`grounded_numbers` , `briefs_cite_contract`). Logged to an MLflow evaluation run for version-over-version comparison. (*First run: relevance/safety/grounded = 1.0, briefs-cite-contract = 0.8, correctness = 0.75.*)
- **Online monitoring** (`monitor_agent.py`). Scheduled scorers (Safety, RelevanceToQuery, `grounded_numbers`) registered on the agent's MLflow experiment so **sampled production traffic is judged continuously**.
- **Re-register.** When eval regresses, `deploy_agent.py` registers the next model version and redeploys — closing the loop.

Note: AI Gateway *usage tracking* isn't supported for the agent endpoint type in this workspace; the inference table provides the equivalent observability.

4.12 The app — `app.py`, `server/routes/`, `frontend/`

- **Backend** (FastAPI): serves `/api/*` and the built React SPA. Routers: `catalog` (`hub/scorecard/overview/health`), `agentic` (`brief/deck/rehearse` + saved artifacts), `genie_routes` (`ask/scorecard/dashboard`), `supervisor_routes` (the router). The lifespan opens/closes the Lakebase pool.
 - **Frontend** (React + Tailwind, dark teal/orange theme): **Deck Hub** (supplier gallery), **Negotiator** (supervisor chat), **Scorecard** (Genie table + Ask + embedded AI/BI dashboard), **Brief** (with clause-citation chips), **Deck**, **Rehearsal**. The sidebar shows a live **LLM-routing badge**.
-

5. End-to-end request flows

A negotiator asks the Negotiator: "draft a brief for Pepsi to claw back COGS"

```
Browser → /api/supervisor/ask
  → supervisor classifies: route=brief, supplier=pepsi
  → agents.build_negotiation_brief(pepsi, objective)
    → knowledge.retrieve(...) → Vector Search index → top contract clauses
    → ChatPromptTemplate | get_llm() → Claude (Mosaic) → grounded brief
  → state.save_artifact("briefs", ...) → Lakebase
  → returns {route:"brief", answer, sources} → UI renders brief + clause chips
```

The Scorecard table loads

```
Browser → /api/genie/scorecard
  → genie.ask(fixed question) → Genie Conversation API
    → Genie writes SQL over kroger_demo.cogs → runs on the warehouse
  → returns {text, sql, columns, rows} → UI renders table + "Show SQL"
```

Someone queries the served model directly

```
POST /serving-endpoints/agents_.../invocations
  {messages, custom_inputs:{task:"brief", supplier_key:"kdp"}}
  → ChatAgent.predict → build_negotiation_brief → Vector Search + get_llm() → brief
  → inference tables capture the request/response
```

6. Deploy & run

Prerequisites

- Databricks CLI authenticated to the workspace (profile `cogs-demo`)
- `uv` (Python) and `node` / `npm` (frontend)
- Serverless workspace (for Lakebase + Apps)

One-time data/AI build (workspace-side)

```
DATABRICKS_CONFIG_PROFILE=cogs-demo python data_foundation.py      # Delta + Metric View
DATABRICKS_CONFIG_PROFILE=cogs-demo python build_genie.py         # Genie space
DATABRICKS_CONFIG_PROFILE=cogs-demo python build_knowledge.py     # Volume + Vector Search
DATABRICKS_CONFIG_PROFILE=cogs-demo python build_dashboard.py    # AI/BI dashboard
```

The app

```
# build frontend
cd frontend && npm install && npm run build && cd ..
# lock deps
uv lock --native-tls
# sync + deploy
databricks sync . /Users/<you>/cogs-negotiation-agent --exclude node_modules --exclude .v
databricks apps deploy cogs-negotiation-agent \
  --source-code-path /Workspace/Users/<you>/cogs-negotiation-agent -p cogs-demo
```

Attach resources (UI or CLI): the **LLM serving endpoint** (`CAN_QUERY`) and the **Lakebase database** (`CAN_CONNECT_AND_CREATE`).

The served model

```
# needs extra local tooling (not app runtime deps):
uv pip install mlflow databricks-agents azure-storage-file-datalake azure-storage-blob az
DATABRICKS_CONFIG_PROFILE=cogs-demo python deploy_agent.py
```

Local dev

```
export DATABRICKS_PROFILE=cogs-demo
uv run uvicorn app:app --reload --port 8000      # backend
cd frontend && npm run dev                        # frontend (proxies /api → :8000)
```

7. Switching the LLM provider (Mosaic ↔ LiteLLM)

In `app.yaml` (app) or the serving endpoint env (served model):

```
# Mosaic (default, governed)
LLM_PROVIDER: mosaic
LLM_MODEL:    databricks-claude-sonnet-4-5

# LiteLLM
LLM_PROVIDER: litellm
LITELLM_BASE_URL: http://<litellm-proxy>:4000
LITELLM_API_KEY:  {{secrets/<scope>/<key>}}
LITELLM_MODEL:    claude-sonnet-4
LITELLM_ROUTES_VIA_MOSAIC: "true"    # keep one governance boundary
```

Then redeploy. No agent code changes.

8. Design decisions & known limitations

- **Synthetic-but-real data.** Numbers originate in `server/data.py` and are materialized as real Delta tables; swap ingestion + a certified metric layer for production.
- **Served model is Genie-free** (self-contained); the app handles Genie NL→SQL where the app SP has working data access.
- **Regional viz uses bars, not a choropleth** — the data is by Census division, not US states/countries (a choropleth needs standard region keys).
- **"Avg COGS Inflation"** on the dashboard is a simple average (filter-friendly); the spend-weighted "blended" figure lives in the metric view for exact reporting.
- **Embedding** is enabled with a least-privilege workspace policy (`ALLOW_APPROVED_DOMAINS`, only the app's domain).

9. Resource inventory

Resource	Identifier
Workspace	<code>adb-7405614449041750.10.azure.databricks.net</code> (Azure) · CLI profile <code>cogs-demo</code>
App URL	<code>https://cogs-negotiation-agent-7405614449041750.10.azure.databricksapps.com</code>

Resource	Identifier
App service principal	a432c266-0858-4f97-aa29-4c726a30c0eb
Catalog / schema	kroger_demo.cogs
Delta tables	supplier_scorecard , regional_performance , contract_chunks
Metric View	kroger_demo.cogs.landed_cost_metrics
Volume	/Volumes/kroger_demo/cogs/contracts/
Genie space	01f1642219cf135cb84f7e0dbc8d6957
Vector Search index	kroger_demo.cogs.contract_chunks_index (endpoint kroger-recipe-search , embeddings databricks-gte-large-en)
Registered model	kroger_demo.models.cogs_negotiation_agent
Agent serving endpoint	agents_kroger_demo-models-cogs_negotiation_agent
Foundation model	databricks-claude-sonnet-4-5
Lakebase instance	cogs-lakebase (ep-blue-pine-e121z0sk.database.eastus2.azuredatabricks.net)
AI/BI dashboard	01f1643845191b049c9a3acf7531495b
SQL warehouse	a455a68035c1f578 (Serverless Starter)